

# Generative Models

*IEOR 135 / 290: Applied Data Science with Venture Applications*

**Alexander Fred-Ojala**  
Research Director, Data Lab  
SCET, UC Berkeley  
[afo@berkeley.edu](mailto:afo@berkeley.edu)



- Generative models
- Deep Dream
- Neural Style Transfer
- PixelRNNs / PixelCNNs
- Variational Autoencoders (+ Autoencoders)
- Generative Adverserial Networks (GANs)
- Examples + Notebook

*Let the computer be an artist!*

# **Supervised vs. Unsupervised**

# **Discriminative vs. Generative**

# Supervised vs. Unsupervised learning

## Supervised learning

Learn function to map input  $x$  to output  $y$ .

## Unsupervised learning

No labels - utilize more data!

Goal: Learn underlying representation of the data.

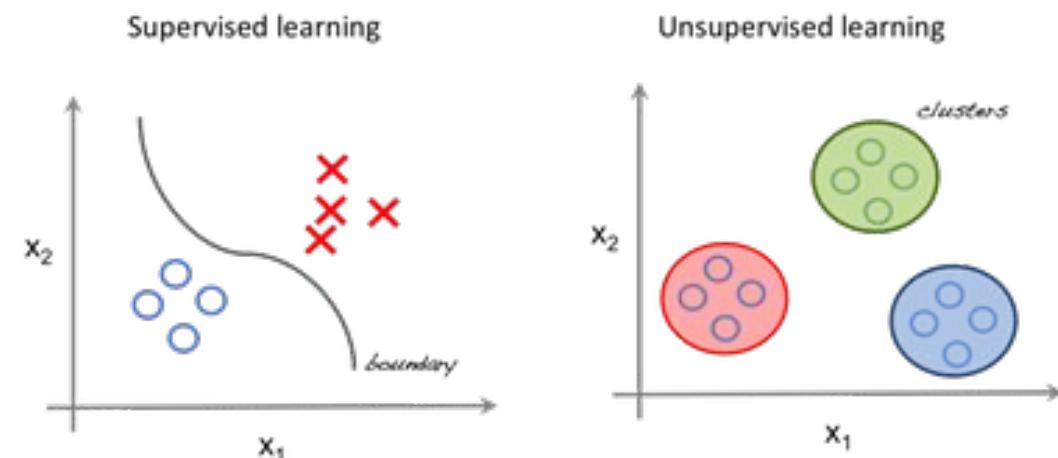


Image source: <https://towardsdatascience.com/unsupervised-learning-with-python-173c51dc7f03>

# Supervised vs. Unsupervised learning

## Supervised learning

**Data:**  $(x, y)$

$x$  is data,  $y$  is label

**Goal:** Learn a *function* to map  $x \rightarrow y$

**Examples:** Classification,  
regression, object detection,  
semantic segmentation, image  
captioning, etc.



→ Cat

Classification

Source: Serena Yeung, CS231n, Stanford

# Supervised vs. Unsupervised learning

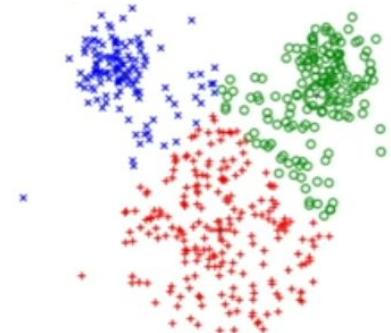
## Unsupervised learning

**Data:**  $x$

Just data, no labels!

**Goal:** Learn some underlying hidden *structure* of the data

**Examples:** Clustering, dimensionality reduction, feature learning, density estimation, etc.

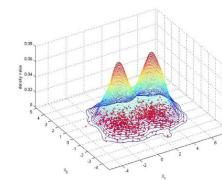


K-means clustering

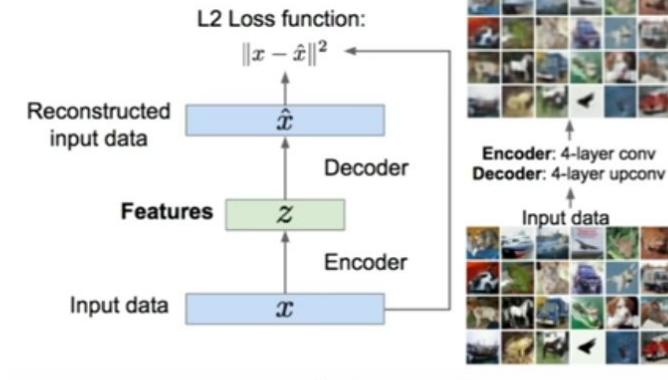


Figure copyright Ian Goodfellow, 2016. Reproduced with permission.

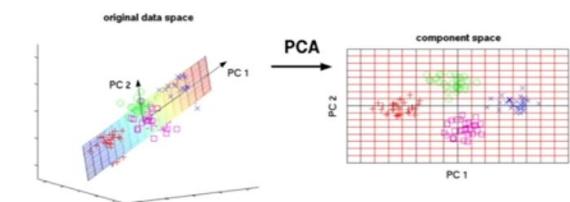
1-d density estimation



2-d density estimation



Autoencoders



Principal Component Analysis

# Generative vs. Discriminative Algorithms

---

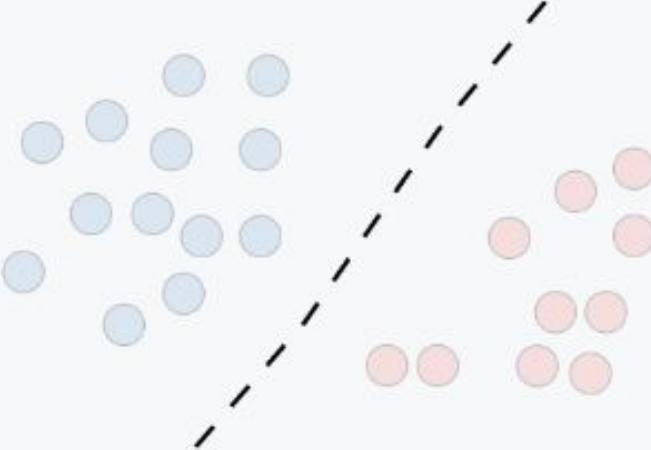
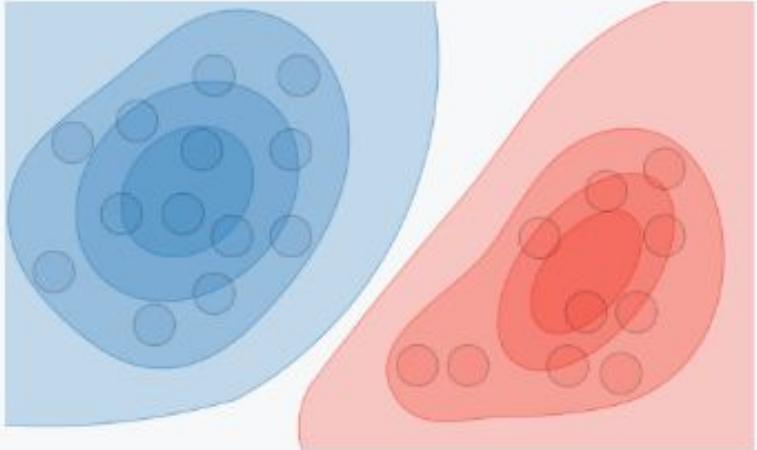
**Discriminative algorithms** classify input data / learn the boundary between classes. Given input  $x$  they predict output  $y$ . Model for classifying emails ( $x$ ) as spam / not spam ( $y$ ):

$p(y|x)$  “*the probability that an email is spam given the words it contains.*”

**Generative algorithms** model the distribution of individual classes / attempt to predict features given a certain label. Modeling the probability of  $x$  given  $y$ :

$p(x|y)$  “*assuming this email is spam, how likely are these words?*”

# Generative vs. Discriminative Algorithms

	Discriminative model	Generative model
Goal	Directly estimate $P(y x)$	Estimate $P(x y)$
What's learned	Decision boundary	Probability distributions of the data
Illustration	 A scatter plot with two classes of data points: blue circles on the left and red circles on the right. A dashed diagonal line, representing the decision boundary, separates the two classes.	 Two overlapping probability density plots. The blue plot is centered on the left side of the slide, and the red plot is centered on the right side. Both plots have a bell-shaped curve with several smaller circles inside, representing individual data points.

# Generative Models

Given training data, generate new samples from same distribution



Training data  $\sim p_{\text{data}}(x)$



Generated samples  $\sim p_{\text{model}}(x)$

Want to learn  $p_{\text{model}}(x)$  similar to  $p_{\text{data}}(x)$

Addresses density estimation, a core problem in unsupervised learning

## Several flavors:

- Explicit density estimation: explicitly define and solve for  $p_{\text{model}}(x)$
- Implicit density estimation: learn model that can sample from  $p_{\text{model}}(x)$  w/o explicitly defining it

# Why Generative Models?

- Realistic samples for artwork, super-resolution, colorization, etc.



- Generative models of time-series data can be used for simulation and planning (reinforcement learning applications!)

## Models that are able to produce / generate new content:

- Generate text
- Generate images
- Generate audio
- Generate number
- Generate distributions
- Generate music notes
- Generate movement

.....

SYSTEM PROMPT  
(HUMAN-WRITTEN)

In a shocking finding, scientist discovered a herd of unicorns living in a remote, previously unexplored valley, in the Andes Mountains. Even more surprising to the researchers was the fact that the unicorns spoke perfect English.

MODEL COMPLETION  
(MACHINE-WRITTEN, 10 TRIES)

The scientist named the population, after their distinctive horn, Ovid's Unicorn. These four-horned, silver-white unicorns were previously unknown to science.

Now, after almost two centuries, the mystery of what sparked this odd phenomenon is finally solved.

Dr. Jorge Pérez, an evolutionary biologist from the University of La Paz, and several companions, were exploring the Andes Mountains when they found a small valley, with no other animals or humans. Pérez noticed that the valley had what appeared to be a natural fountain, surrounded by two peaks of rock and silver snow.

Example from Open-AIs Transformer model GPT2

# Deep Dream

# Deep Dream (July 2015)

Uses a CNN to find and enhance patterns in images to create dream-like images

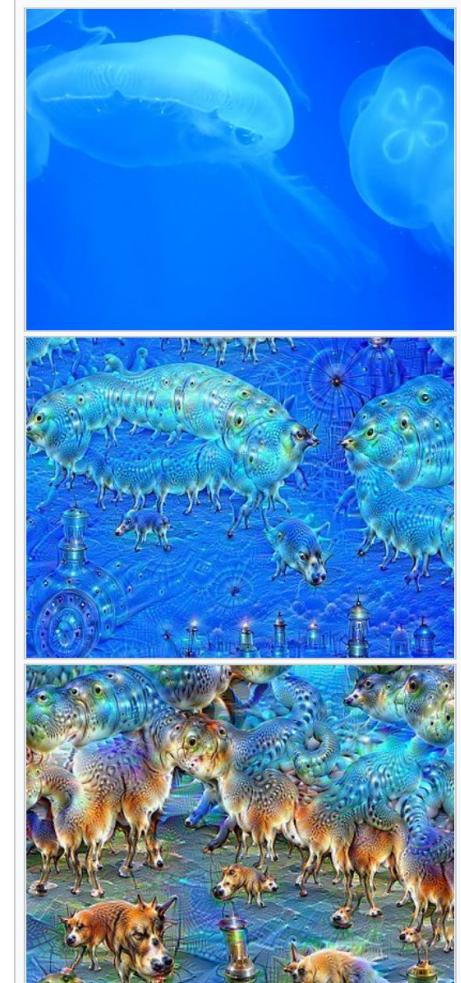


# Deep Dream

## How it works

- Take a pretrained CNN (e.g. VGG16 trained on ImageNet)
- Setup a gradient ascent function, maximizing the value of (arbitrary) activation maps in the bottom layer
- Run the network in reverse, adjusting pixel values in the input image every iteration
- After several iterations psychedelic and surreal images are generated algorithmically.

*It's like backprop, but we don't adjust the network weights. Instead the weights are held fixed and the input is adjusted.*



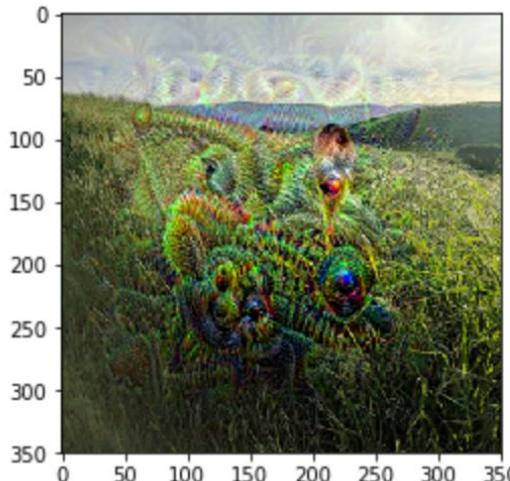
The original image (top) after applying ten (middle) and fifty (bottom) iterations DeepDream, the network having been trained to perceive dogs

# Deep Dream

Example of generating deep dream images:

<https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/8.2-deep-dream.ipynb>

```
from matplotlib import pyplot as plt  
  
plt.imshow(deprocess_image(np.copy(img)))  
plt.show()
```



# Neural Style Transfer

# Neural Style Transfer (Aug 2015)

Use Deep Learning to manipulate digital images / videos to adopt the style of another image.

**Common uses:**

Transfer one image style onto another.



# Neural Style Transfer

## How it works

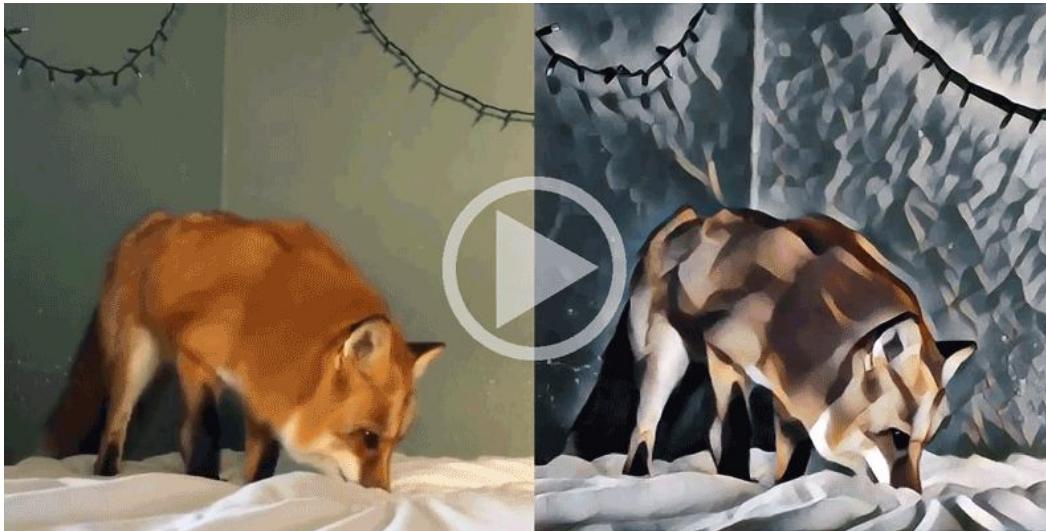


- Take a pretrained CNN (e.g. VGG19 trained on ImageNet)
- Setup a loss function that can capture the content loss (from original image) and the style loss (from style image). Distance can be L2 norm.

```
loss = distance(style(reference_image) - style(generated_image))  
+      distance(content(original_image) - content(generated_image))
```

- Content aims at preserving high-level layer activations between the original image and the generated image.
- Style maintains similar correlations within activations for both low-level layers and high-level layers, between style image and original image.

# Neural Style Transfer



# Neural Style Transfer

Example of using Neural Style Transfer with Tensorflow 2 to create a Kadinsky turtle:

[https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/r2/tutorials/generative/style\\_transfer.ipynb](https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/r2/tutorials/generative/style_transfer.ipynb)

## Take aways

- Style transfer consists in creating a new image that preserves the "contents" of a target image while also capturing the "style" of a reference image.
- "Content" can be captured by the high-level activations of a convnet.
- "Style" can be captured by the internal correlations of the activations of different layers of a convnet.
- Hence deep learning allows style transfer to be formulated as an optimization process using a loss defined with a pre-trained convnet.
- Starting from this basic idea, many variants and refinements are possible!



Image of Green Sea Turtle -By P.Lindgren CC BY-SA 3.0, from Wikimedia Commons



Now how would it look like if Kandinsky decided to paint the picture of this Turtle exch



# PixelRNN / PixelCNN

## Fully visible belief network

Explicit density model

Use chain rule to decompose likelihood of an image  $x$  into product of 1-d distributions:

$$p(x) = \prod_{i=1}^n p(x_i | x_1, \dots, x_{i-1})$$

↑                              ↑  
Likelihood of                  Probability of i'th pixel value  
image  $x$                           given all previous pixels

Then maximize likelihood of training data

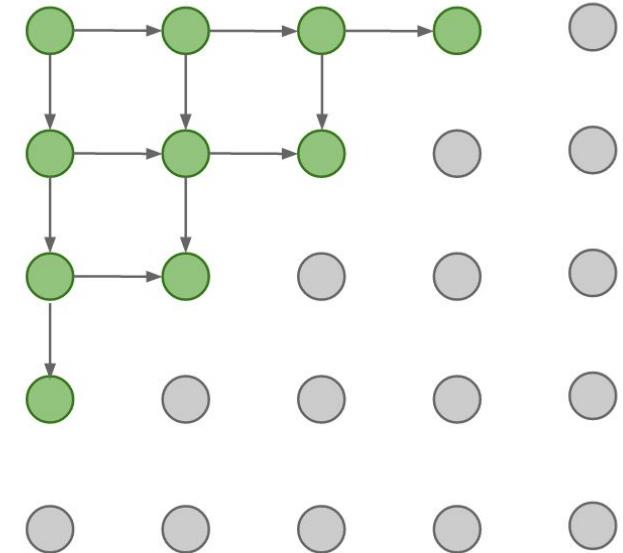
Source: Serena Yeung, CS231n, Stanford

# Pixel RNNs

Generate image pixels starting from corner

Dependency on previous pixels modeled  
using an RNN (LSTM)

Drawback: sequential generation is slow!



Source: Serena Yeung, CS231n, Stanford

# Pixel CNNs

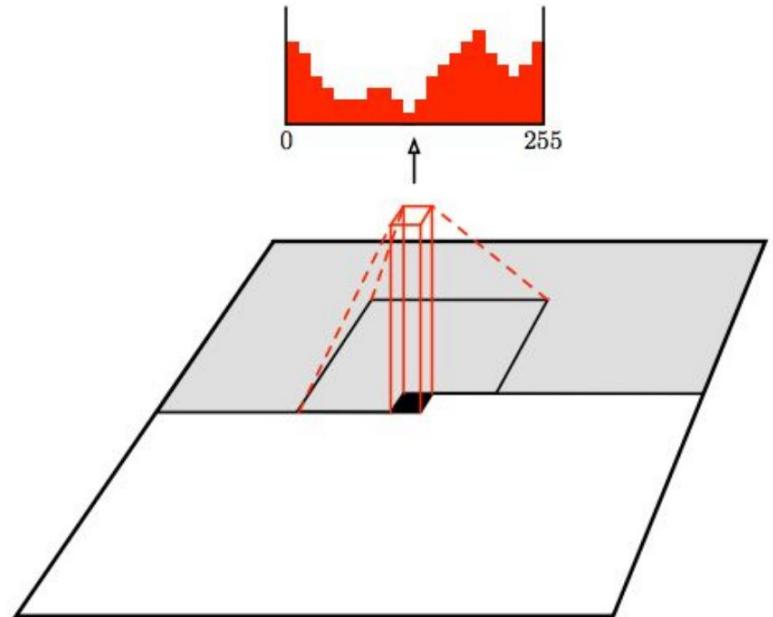
Still generate image pixels starting from corner

Dependency on previous pixels now modeled using a CNN over context region

Training: maximize likelihood of training images

$$p(x) = \prod_{i=1}^n p(x_i|x_1, \dots, x_{i-1})$$

Softmax loss at each pixel



Training is faster than PixelRNN  
(can parallelize convolutions since context region values known from training images)

Generation must still proceed sequentially  
=> still slow

Source: Serena Yeung, CS231n, Stanford

# Autoencoders

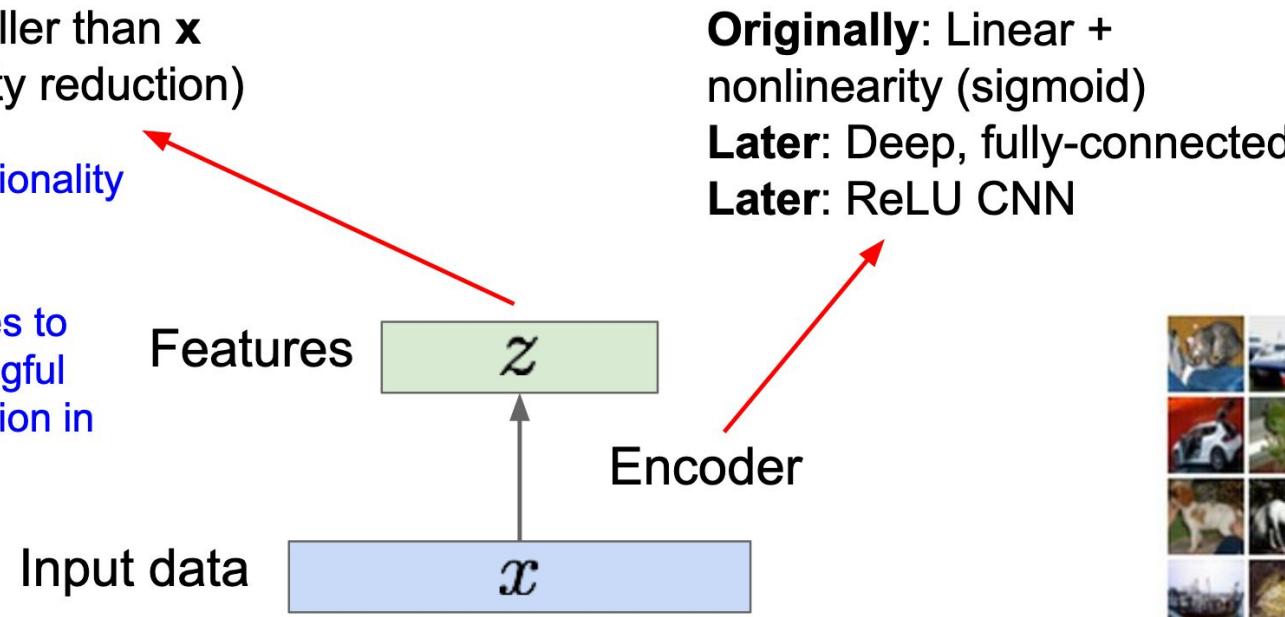
# Autoencoders

Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data

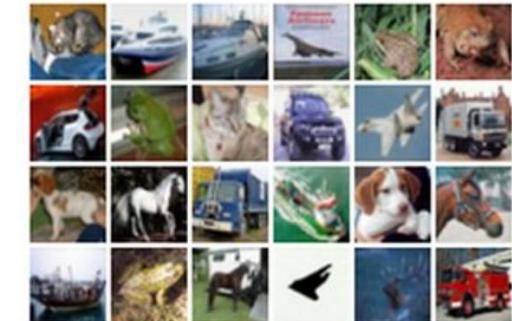
$z$  usually smaller than  $x$   
(dimensionality reduction)

Q: Why dimensionality reduction?

A: Want features to capture meaningful factors of variation in data



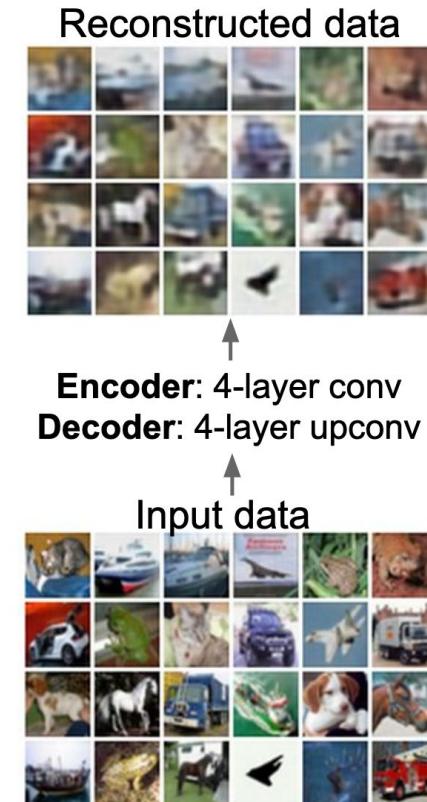
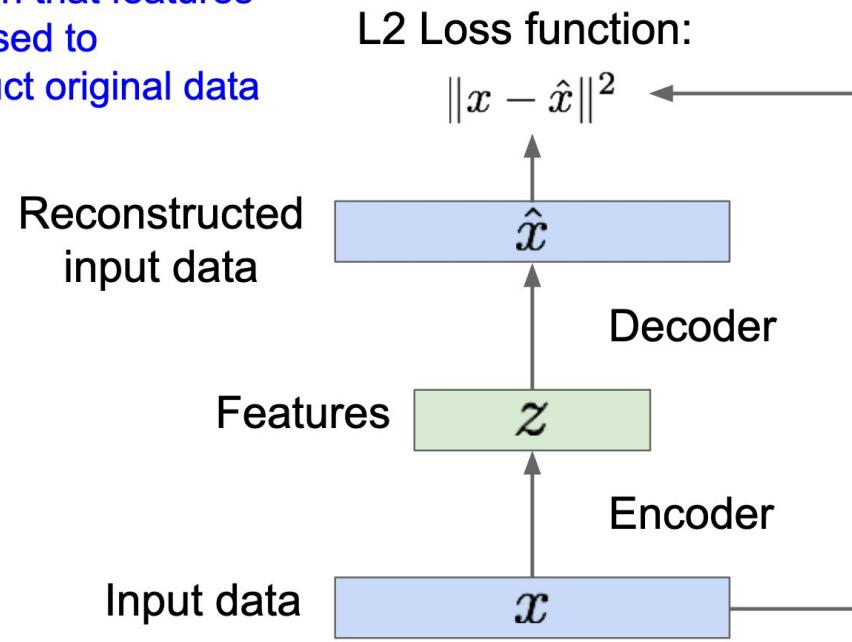
**Originally:** Linear +  
nonlinearity (sigmoid)  
**Later:** Deep, fully-connected  
**Later:** ReLU CNN



Source: Serena Yeung, CS231n, Stanford

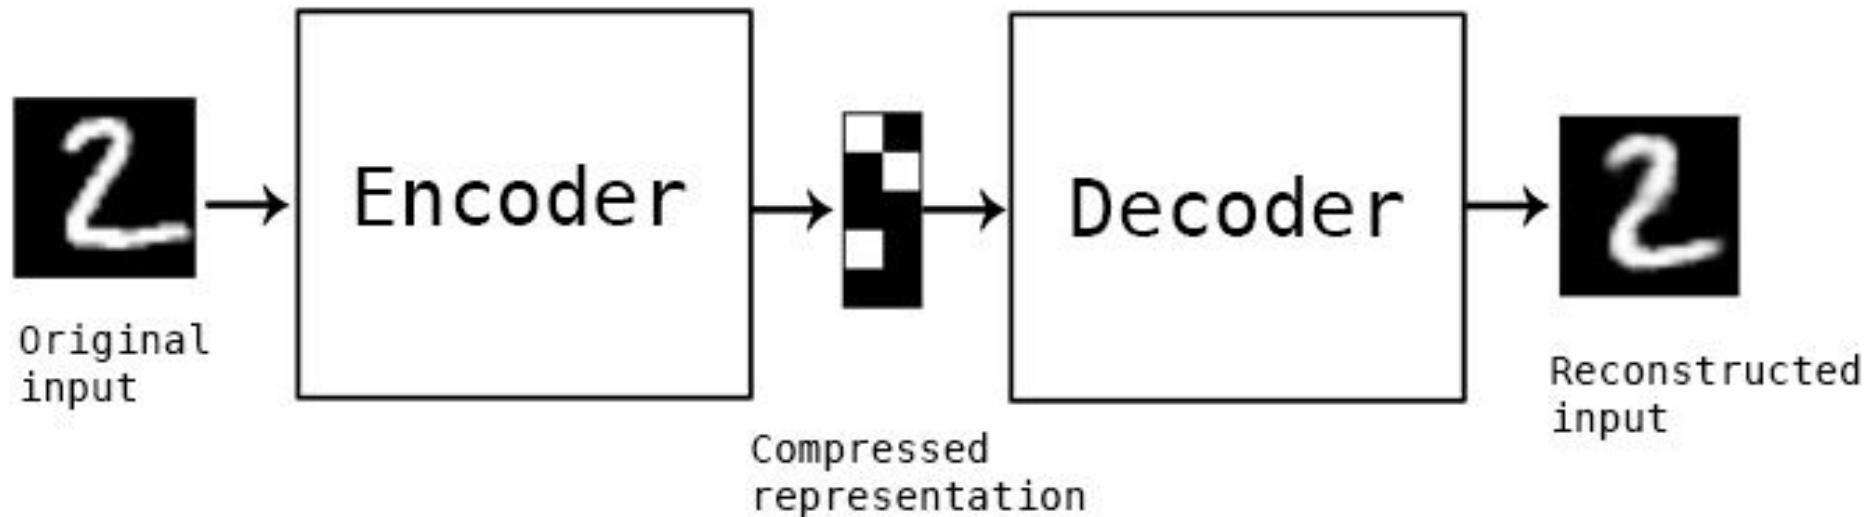
# Autoencoders

Train such that features can be used to reconstruct original data



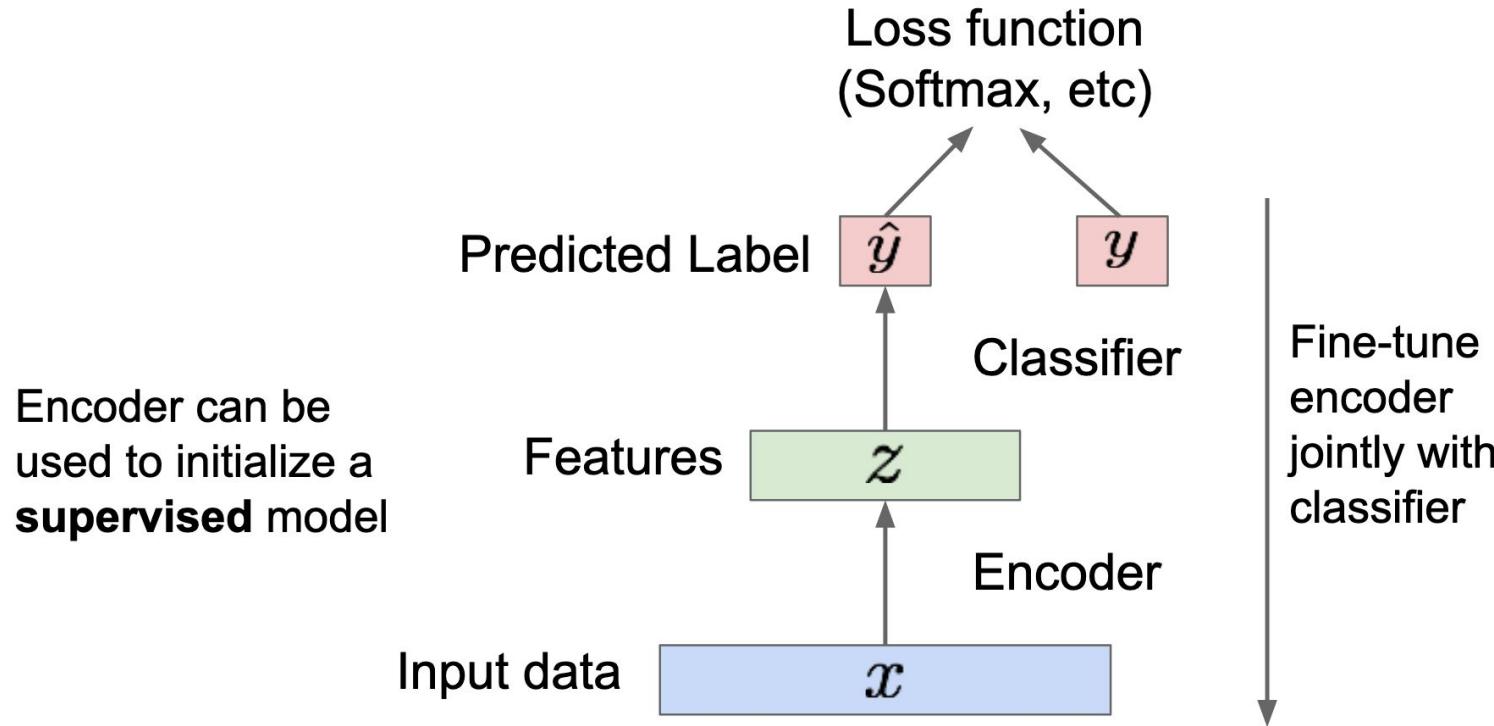
Source: Serena Yeung, CS231n, Stanford

# Autoencoders



# Autoencoders

After training throw away decoder and instead initialize supervised model



Source: Serena Yeung, CS231n, Stanford

# Variational Autoencoders

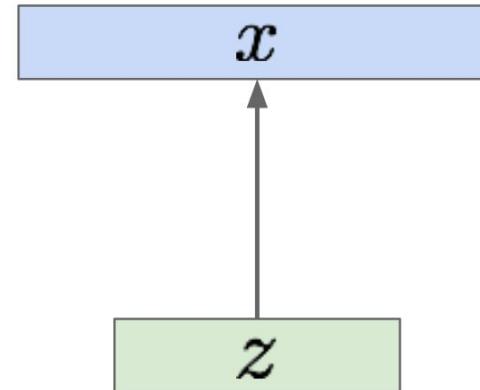
# Variational Autoencoders

Probabilistic spin on autoencoders - will let us sample from the model to generate data!

Assume training data  $\{x^{(i)}\}_{i=1}^N$  is generated from underlying unobserved (latent) representation  $z$

Sample from  
true conditional

$$p_{\theta^*}(x \mid z^{(i)})$$



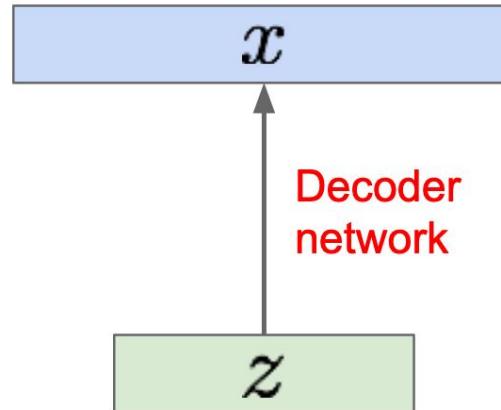
Sample from  
true prior

$$p_{\theta^*}(z)$$

**Intuition** (remember from autoencoders!):  
 $x$  is an image,  $z$  is latent factors used to  
generate  $x$ : attributes, orientation, etc.

# Variational Autoencoders

Sample from  
true conditional  
 $p_{\theta^*}(x \mid z^{(i)})$



Sample from  
true prior  
 $p_{\theta^*}(z)$

We want to estimate the true parameters  $\theta^*$  of this generative model.

How should we represent this model?

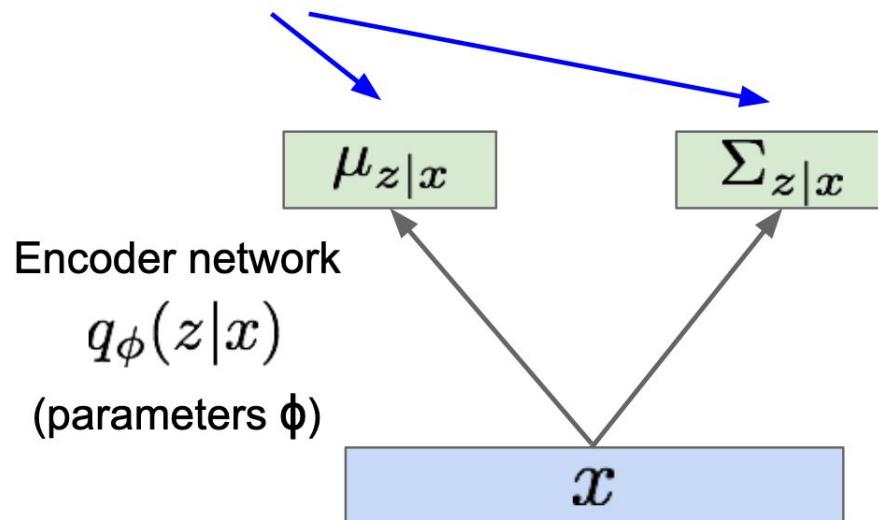
Choose prior  $p(z)$  to be simple, e.g.  
Gaussian.

Conditional  $p(x|z)$  is complex (generates  
image) => represent with neural network

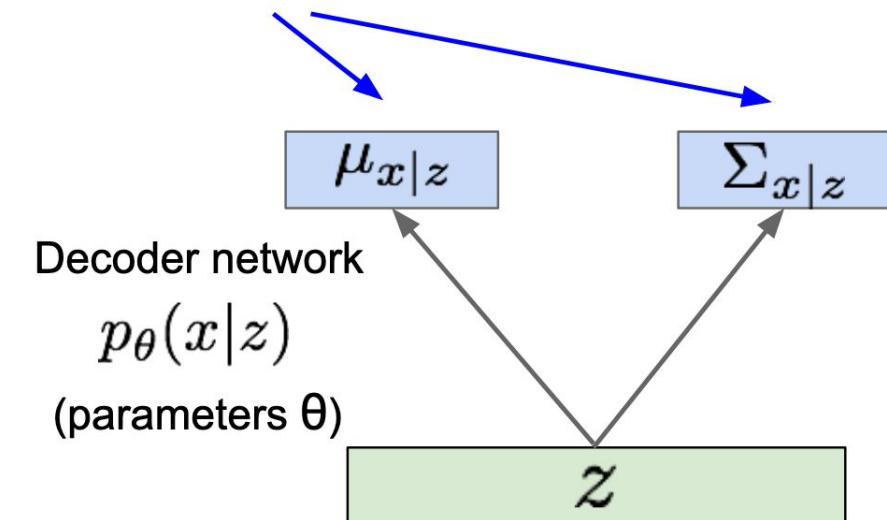
# Variational Autoencoders

Since we're modeling probabilistic generation of data, encoder and decoder networks are probabilistic

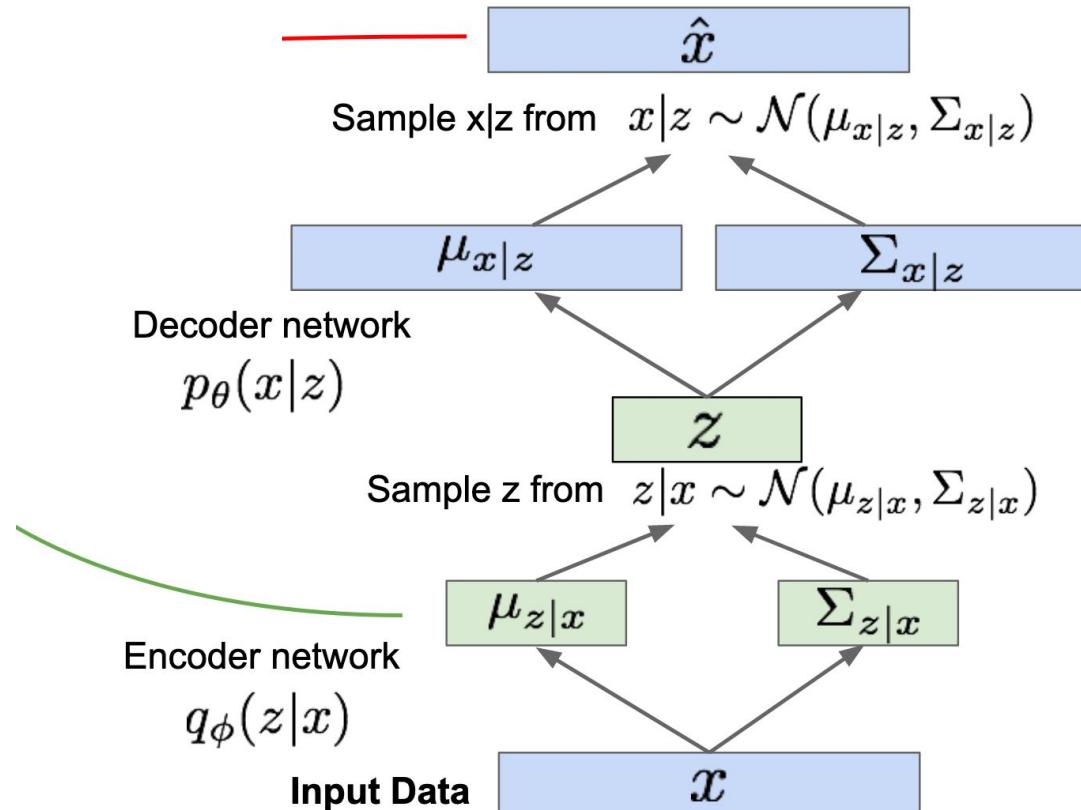
Mean and (diagonal) covariance of  $\mathbf{z} | \mathbf{x}$



Mean and (diagonal) covariance of  $\mathbf{x} | \mathbf{z}$

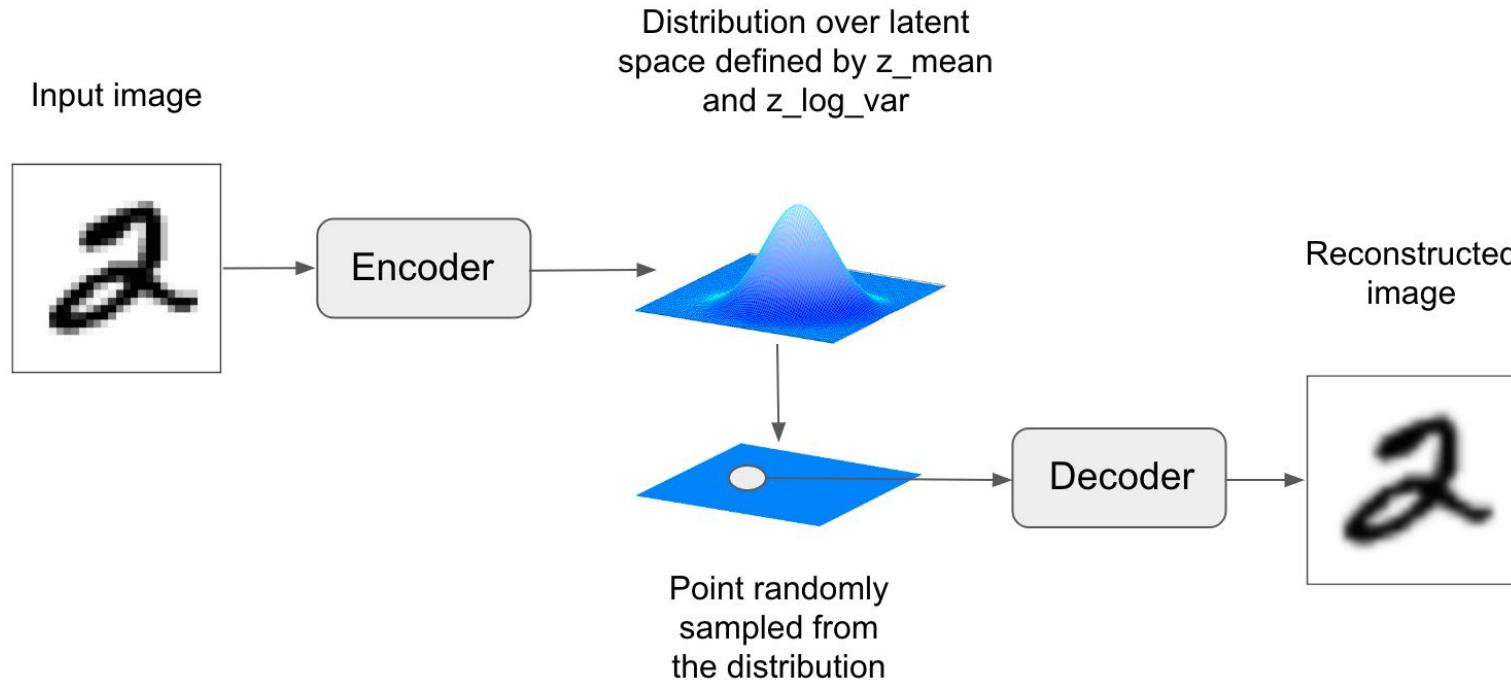


# Variational Autoencoders



Source: Serena Yeung, CS231n, Stanford

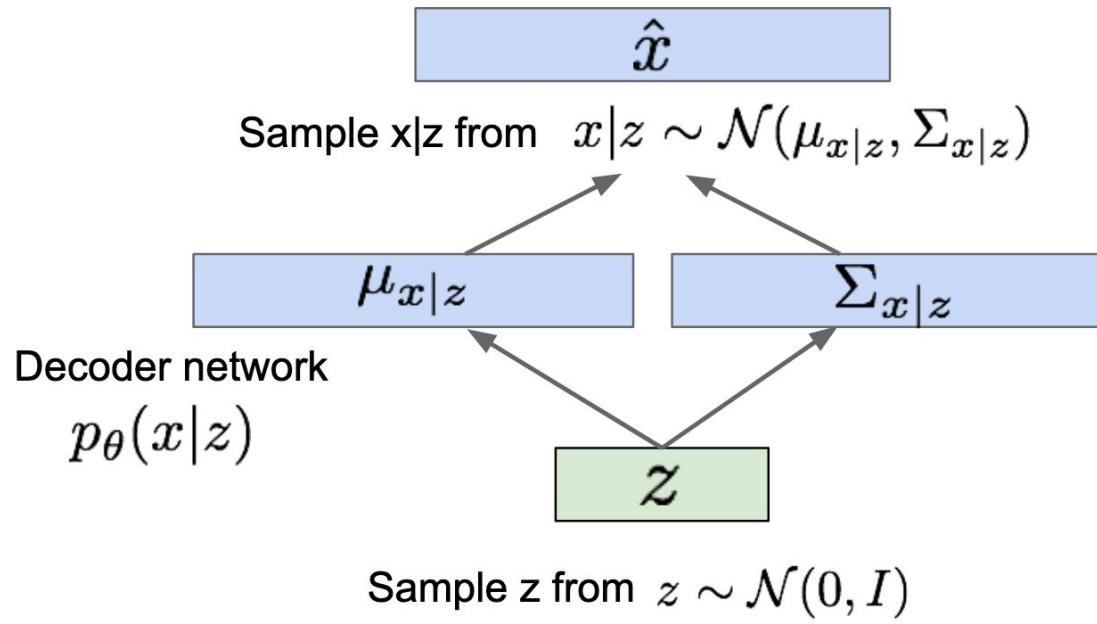
# Variational Autoencoders



Source: keras.io

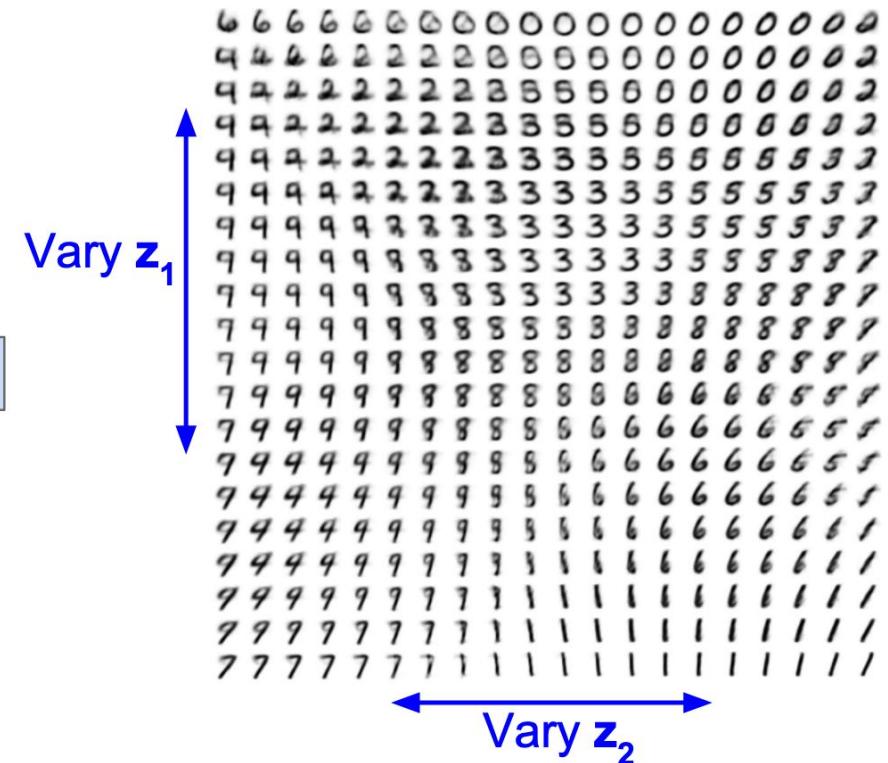
# Variational Autoencoders

Use decoder network. Now sample z from prior!



Kingma and Welling, "Auto-Encoding Variational Bayes", ICLR 2014

## Data manifold for 2-d $\mathbf{z}$



# Variational Autoencoders

Different dimensions of  $z$  encode interpretable factors of variation

Degree of smile

Vary  $z_1$



Vary  $z_2$

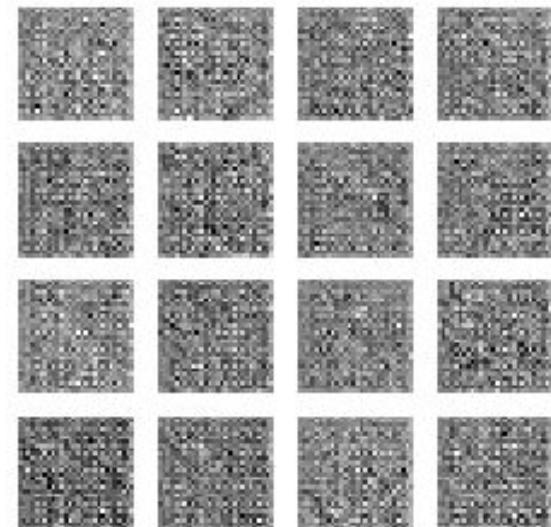
Head pose

Kingma and Welling, "Auto-Encoding Variational Bayes", ICLR 2014

# Convolutional Variational Autoencoder

Generate MNIST images with VAEs in  
Tensorflow 2.0

<https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/r2/tutorials/generative/cvae.ipynb>



# GANs

# GANs History

---

Generative Adversarial Networks (GANs) were developed in 2014 by Ian Goodfellow, Yoshua Bengio et al.

*“The most interesting idea in the last 10 years in Machine Learning”*

- Yann LeCun

Problem: Want to sample from complex, high-dimensional training distribution. No direct way to do this!

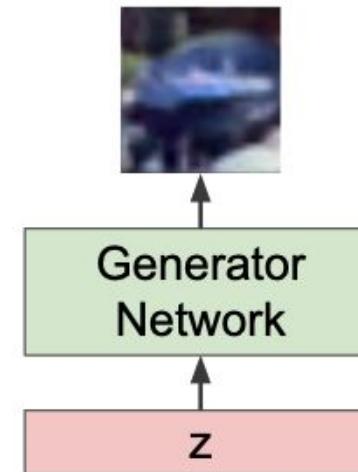
Solution: Sample from a simple distribution, e.g. random noise. Learn transformation to training distribution.

Q: What can we use to represent this complex transformation?

A: A neural network!

Output: Sample from training distribution

Input: Random noise



Source: Serena Yeung, CS231n, Stanford

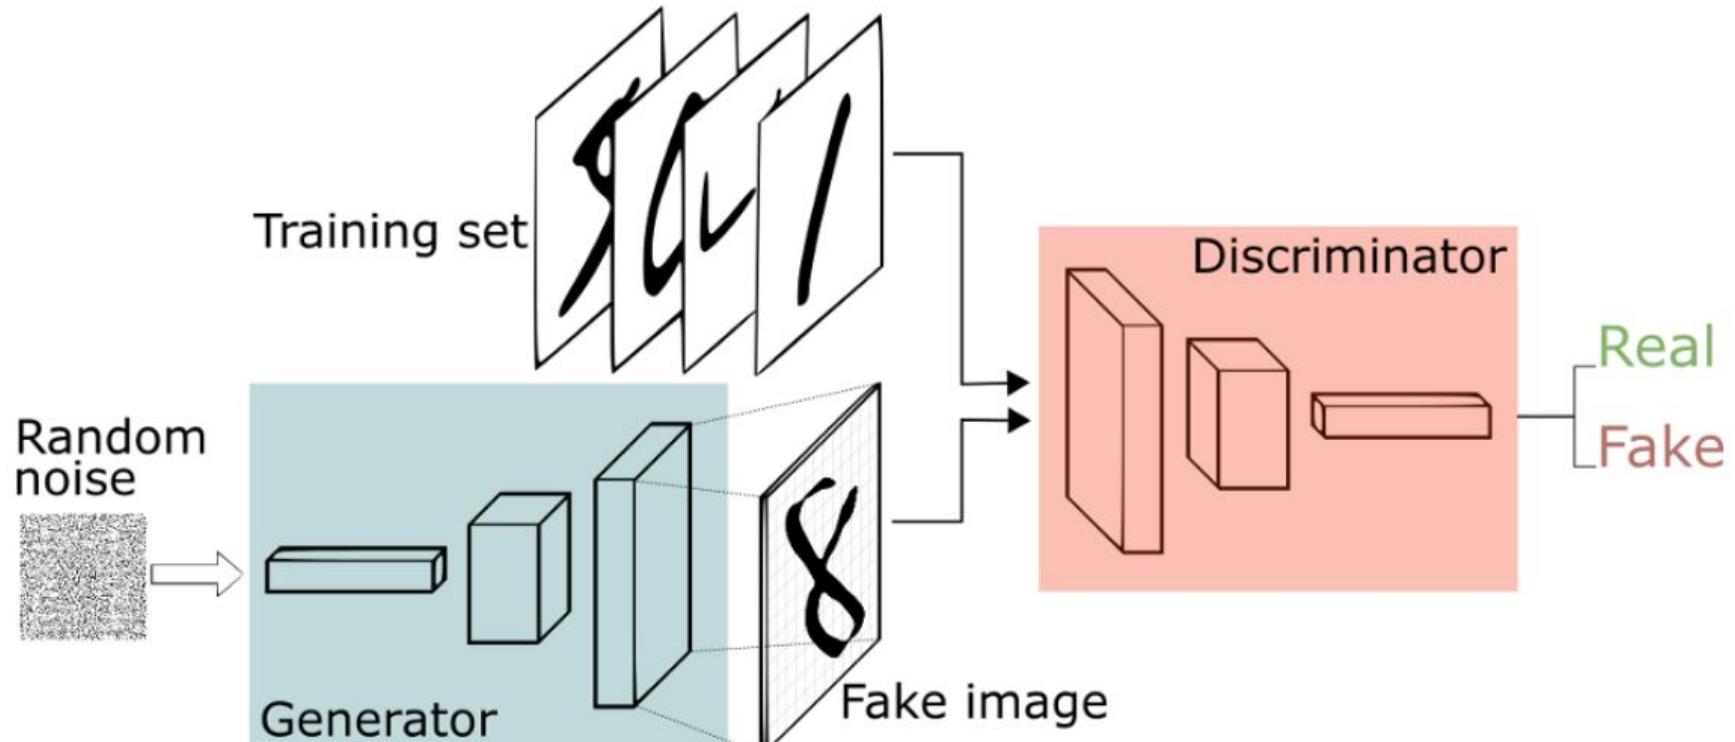
## Consists of Two Functions

*The main idea is to have two separate Neural Networks, a **generator** and a **discriminator**, locked in a competition with each other.*

- ◇ **Generator:** Creates output similar to the one in the dataset
- ◇ **Discriminator:** Tries to understand if they are real or not

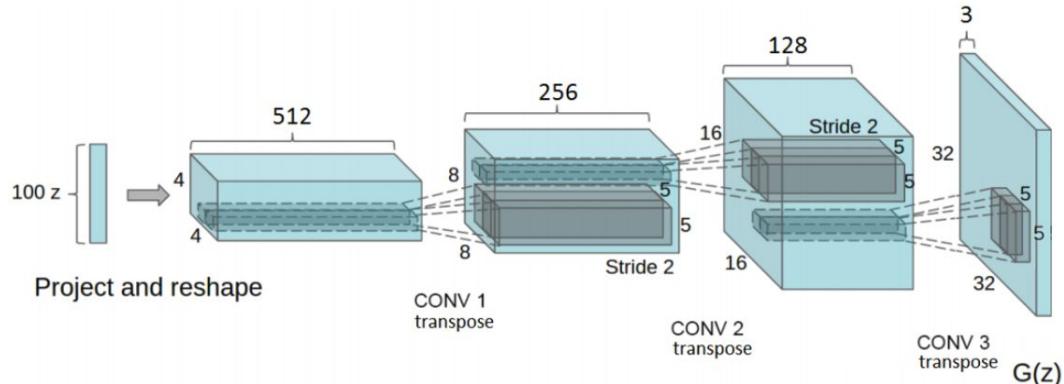
**Metaphor:** The generator function is a money launderer and the discriminator is a police. Cycle of feedback, until the counterfeiter is better than the police.

# GAN Systemic Architecture



Generative Adversarial Network framework.

# GAN Convolutional Generator



Adapted from the DCGAN paper. The Generator network implemented here. Note the non-

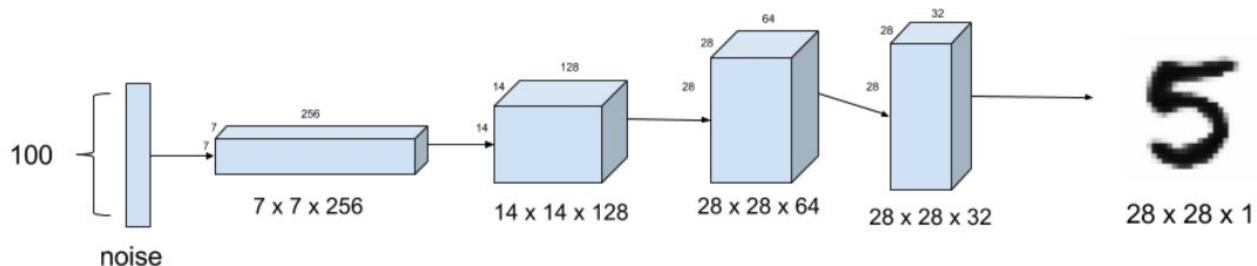
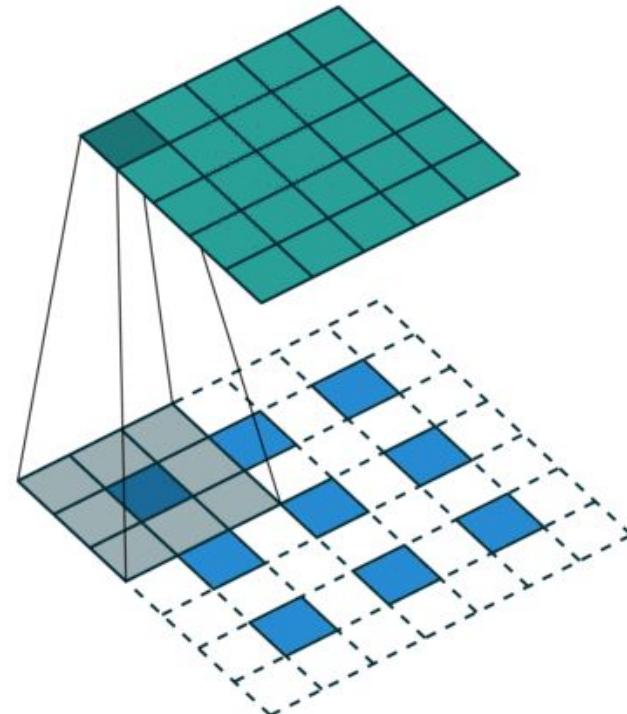
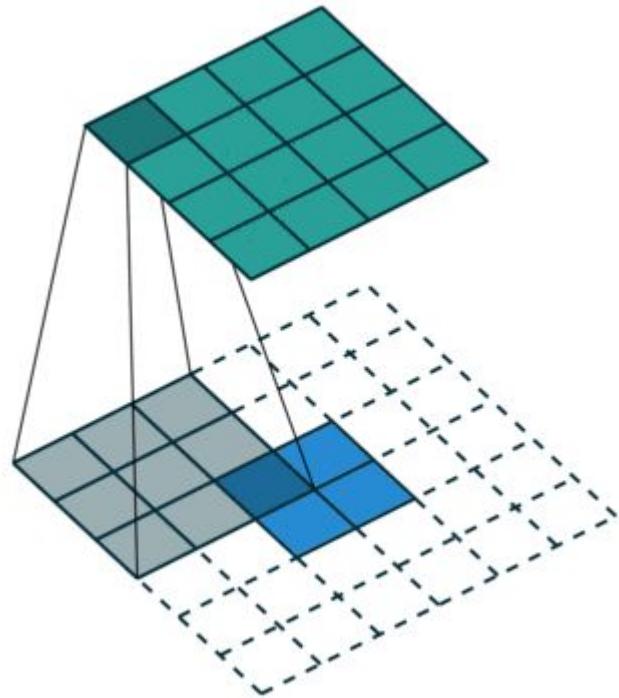


Figure 2. Generator model synthesizes fake MNIST images from noise. Upsampling is used instead of fractionally-strided transposed convolution.

- Source: 1. <https://medium.freecodecamp.org/an-intuitive-introduction-to-generative-adversarial-networks-gans-7a2264a81394>  
2. <https://towardsdatascience.com/gan-by-example-using-keras-on-tensorflow-backend-1a6d515a60d0>

# GAN: Upsampling (transposed convolutions)



Source: 1. <https://medium.freecodecamp.org/an-intuitive-introduction-to-generative-adversarial-networks-gans-7a2264a81394>  
2. <https://towardsdatascience.com/gan-by-example-using-keras-on-tensorflow-backend-1a6d515a60d0>

# GAN: Discriminator

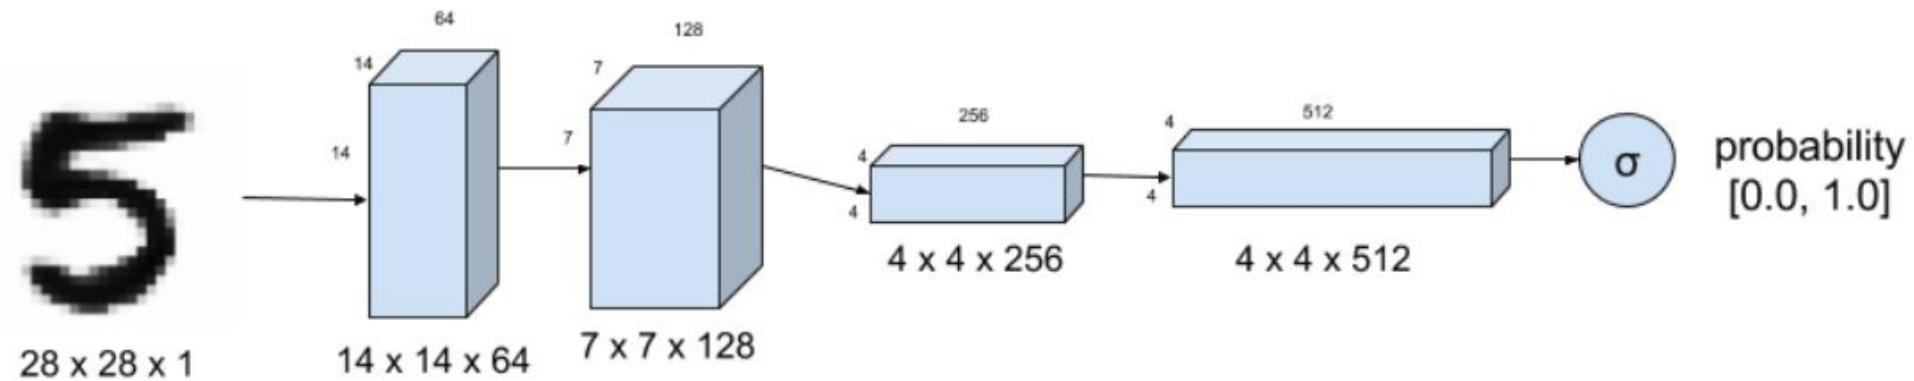


Figure 1. Discriminator of DCGAN tells how real an input image of a digit is. MNIST Dataset is used as ground truth for real images. Strided convolution instead of max-pooling down samples the image.

# GAN Training

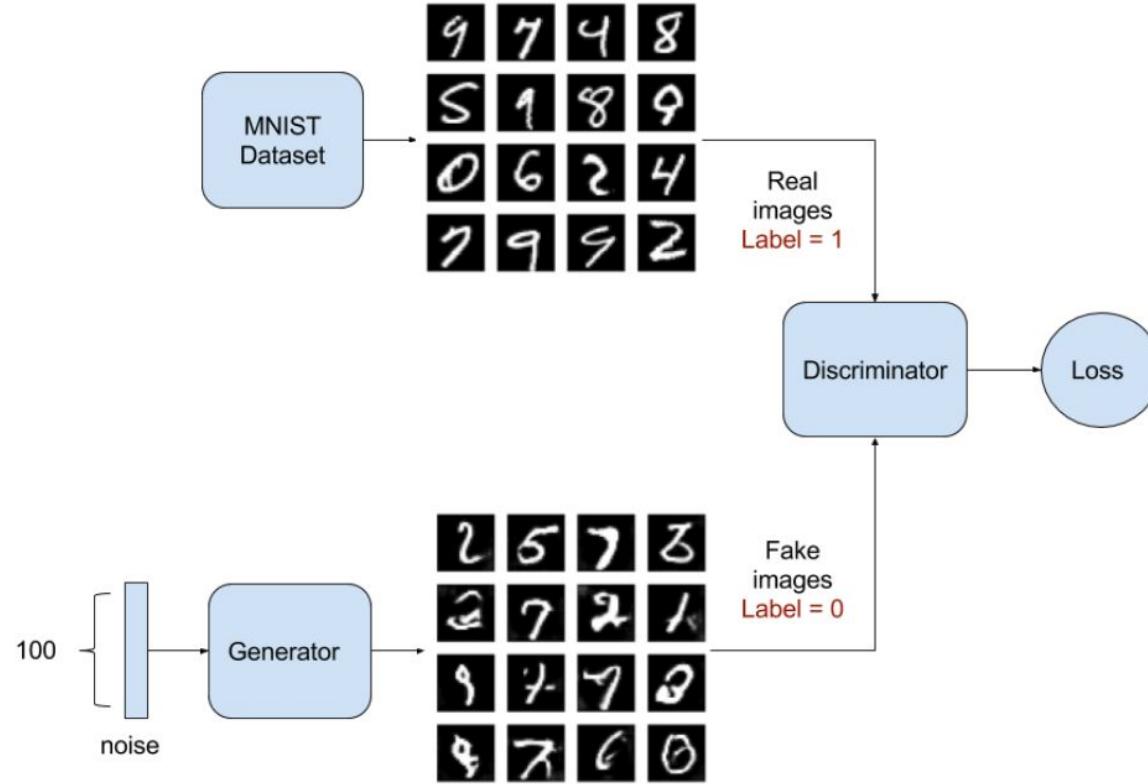


Figure 4. Discriminator model is trained to distinguish real from fake handwritten images.

Source: <https://towardsdatascience.com/gan-by-example-using-keras-on-tensorflow-backend-1a6d515a60d0>

# GAN Training Pseudo code

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{\text{data}}(\mathbf{x})$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D \left( \mathbf{x}^{(i)} \right) + \log \left( 1 - D \left( G \left( \mathbf{z}^{(i)} \right) \right) \right) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log \left( 1 - D \left( G \left( \mathbf{z}^{(i)} \right) \right) \right).$$

**end for**

# The Adversarial Model



Figure 3. The Adversarial model is simply generator with its output connected to the input of the discriminator. Also shown is the training process wherein the Generator labels its fake image output with 1.0 trying to fool the Discriminator.

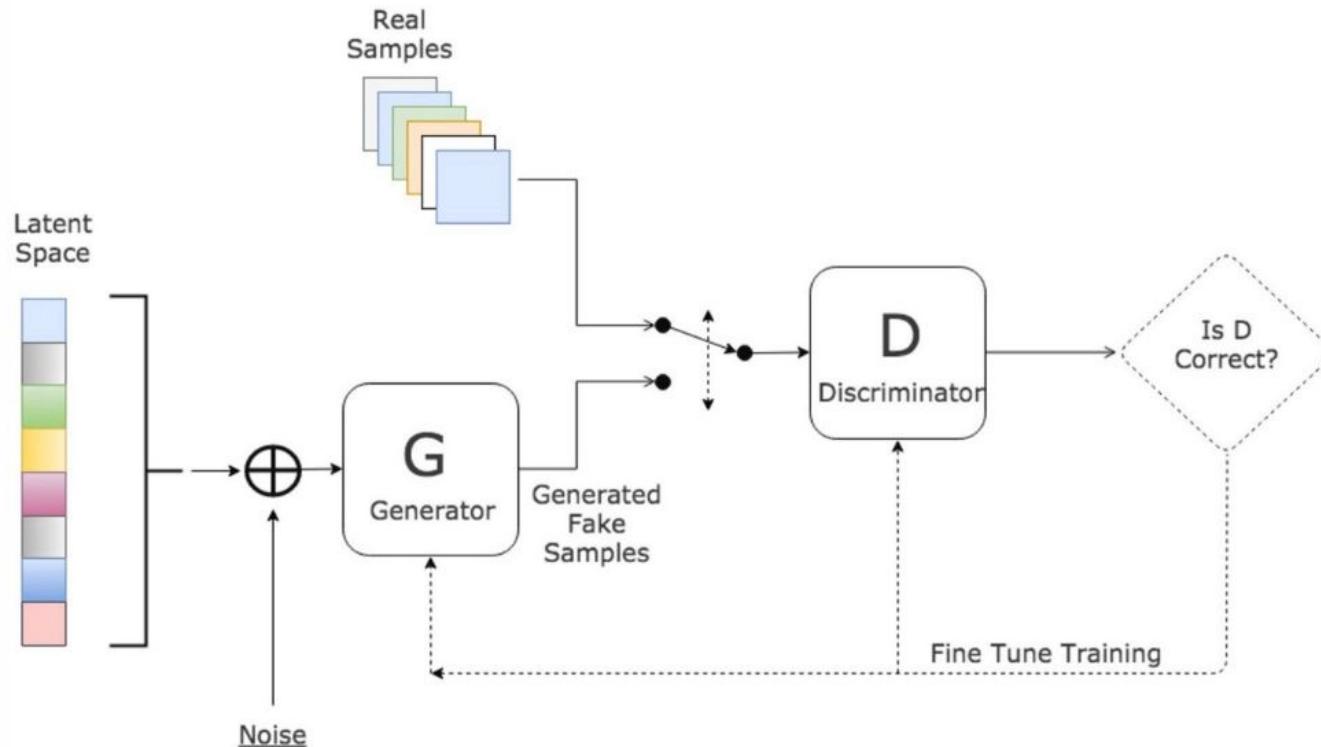
## Train the generator through the adversarial model:

*The Generator is trained indirectly through the adversarial model. Pass noise to the adversarial model and mislabel everything as if they were images taken from the database (even though they are generated). The discriminator, which is trained beforehand on the real images, will fail to label the synthetic images as real, and there will be a high error overall. Then we use backpropagation and the parameters of the discriminator are frozen so only parameters in the generator will be affected. So minimising the error function of the adversarial model really means make the generated images as similar as possible to something, the discriminator will recognise as real.*

# GAN Architecture

## Generative Adversarial Network

*Global concept of a GAN.  
Source Link*



Source: <https://medium.com/ai-society/gans-from-scratch-1-a-deep-introduction-with-code-in-pytorch-and-tensorflow-cb03cdcd8a0f>

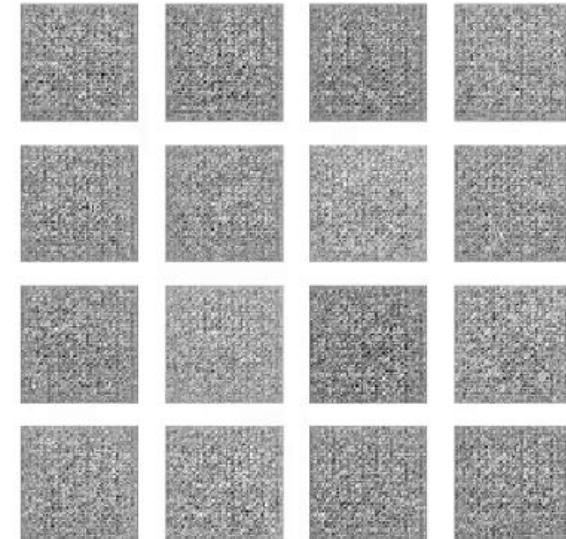
## Result interpretation

*At the **end of training** you want the **loss of the adversarial model to be small, and the error of the discriminator to be high** -- it is no longer able to tell the difference.*

# GAN: Result Examples



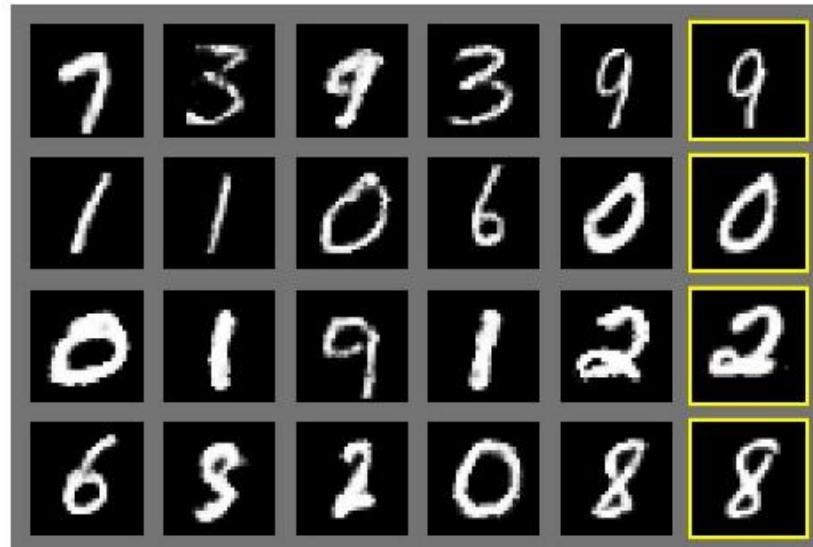
Comparing real (left) and generated (right) MNIST sample images. Because MNIST images have a simpler data structure, the model was able to produce more realistic samples when compared to the SVHNs.



Output of a GAN through time, learning to Create Hand-written digits. We'll code this example!  
<https://medium.com/ai-society/gans-from-scratch-1-a-deep-introduction-with-code-in-pytorch-and-tensorflow-cb03cdcd8a0f>

# GAN: Result Examples

Generated samples



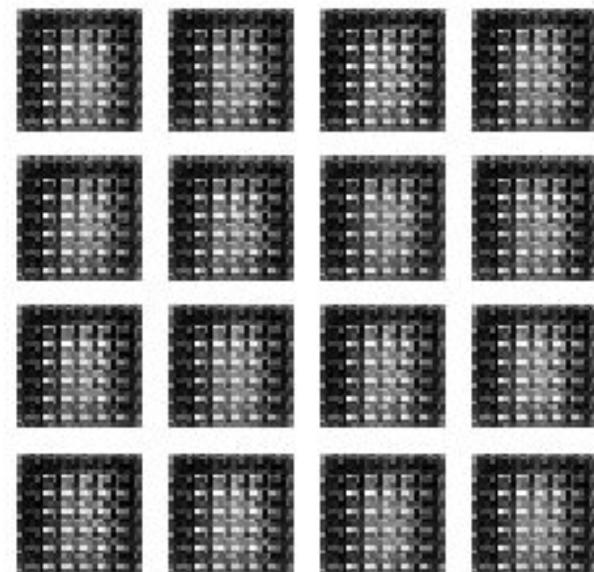
Nearest neighbor from training set

Figures copyright Ian Goodfellow et al., 2014. Reproduced with permission.

# GAN notebook

Example of using DCGANS with  
Tensorflow 2 to generate MNIST data.

<https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/r2/tutorials/generative/dcgan.ipynb>



# GAN Applications

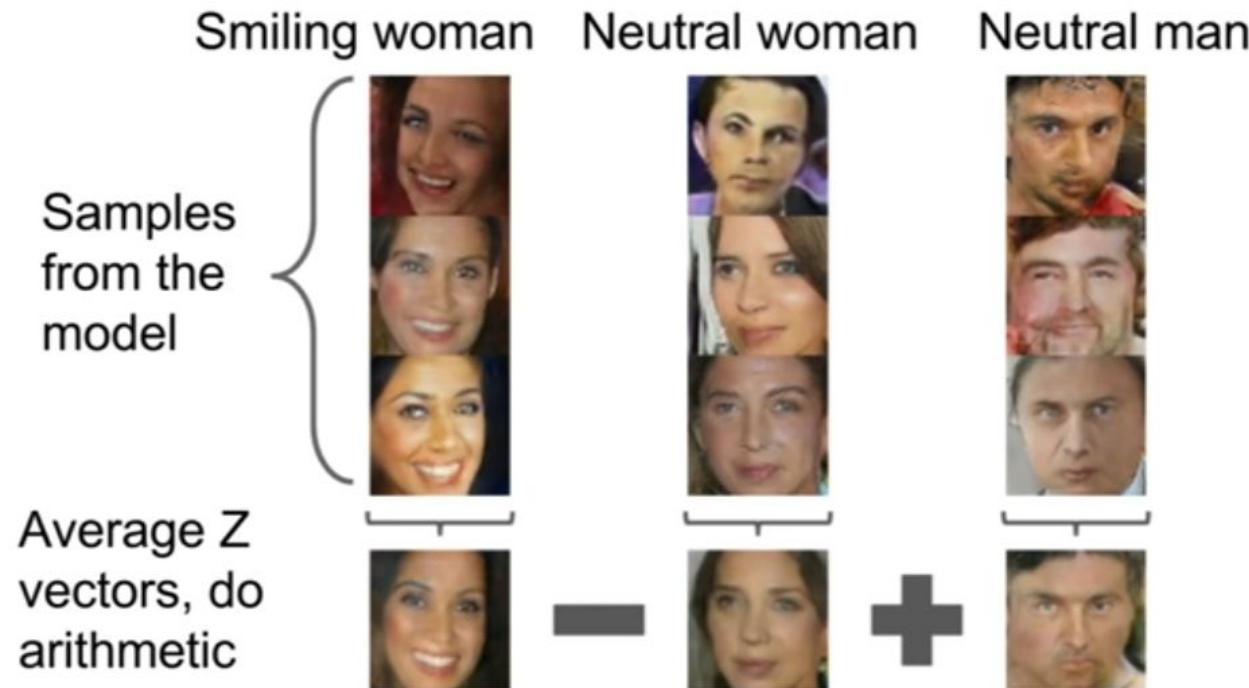
# Tuning the latent trait vector z

*Interpolating between random point (vectors) in the latent space, shows a smooth transition between states.*



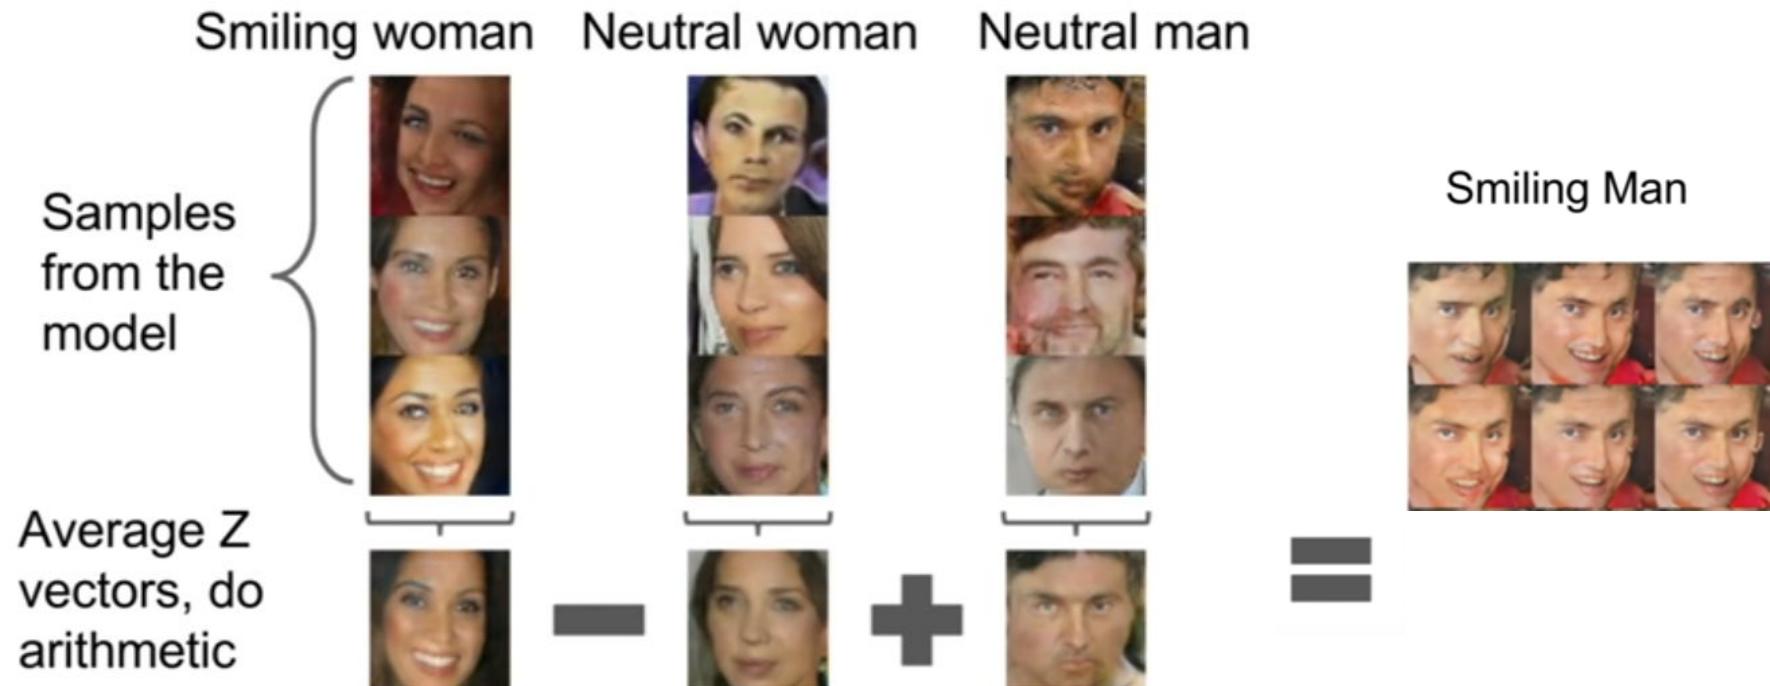
# Interpret GAN generation dimensions

*Finding the latent space dimensions that correspond to certain features in the generated output images. We can try to do vector mathematics of different samples*



# Interpret GAN generation dimensions

*Finding the latent space dimensions that correspond to certain features in the generated output images. We can try to do vector mathematics of different samples*



# Interpret GAN generation dimensions

Glasses man



No glasses man



No glasses woman

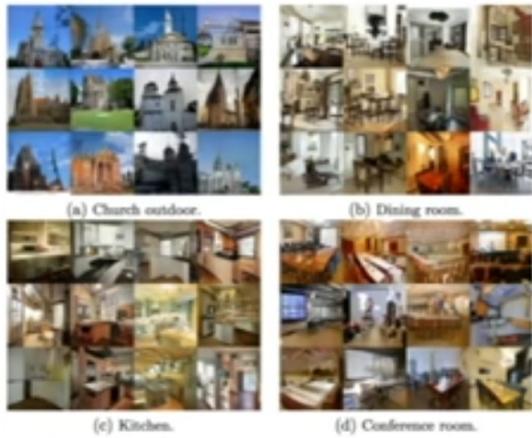


Woman with glasses



# GAN Examples

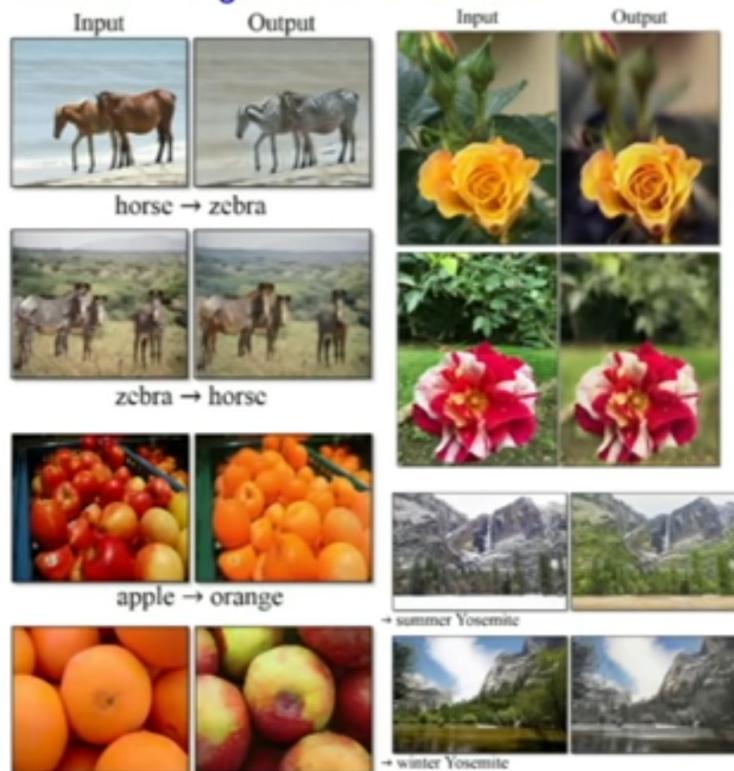
## Better training and generation



LSGAN. Mao et al. 2017.



## Source->Target domain transfer



CycleGAN. Zhu et al. 2017.

## Text -> Image Synthesis

this small bird has a pink breast and crown, and black primaries and secondaries.



this magnificent fellow is almost all black with a red crest, and white cheek patch.



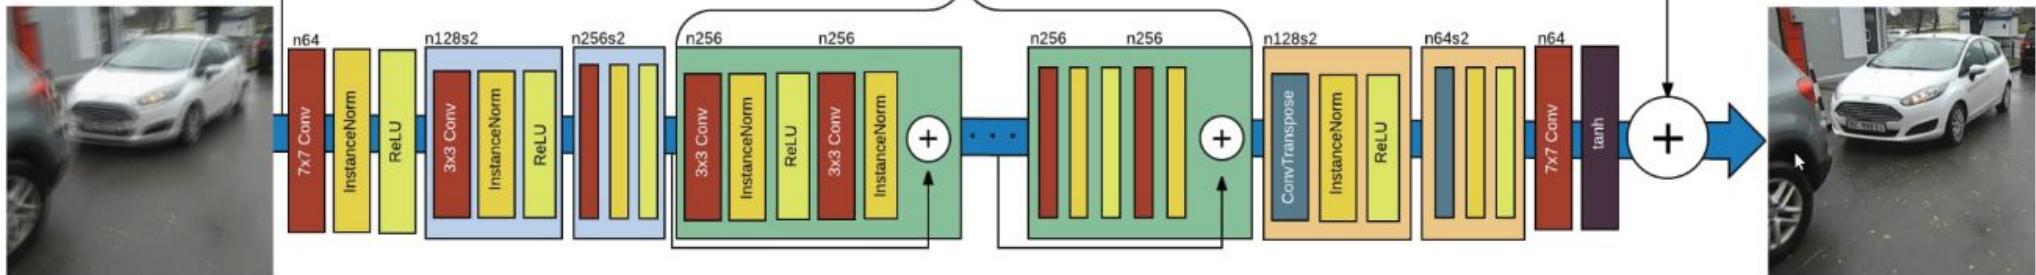
Akata et al. 2017.

## Many GAN applications



Pix2pix.

# GAN Examples: Deblur / Depixel Images



*The Architecture of the DeblurGAN generator network—[Source](#)*

**Source:** <https://blog.sicara.com/keras-generative-adversarial-networks-image-deblurring-45e3ab6977b5>

# Cycle GAN: Style Translation

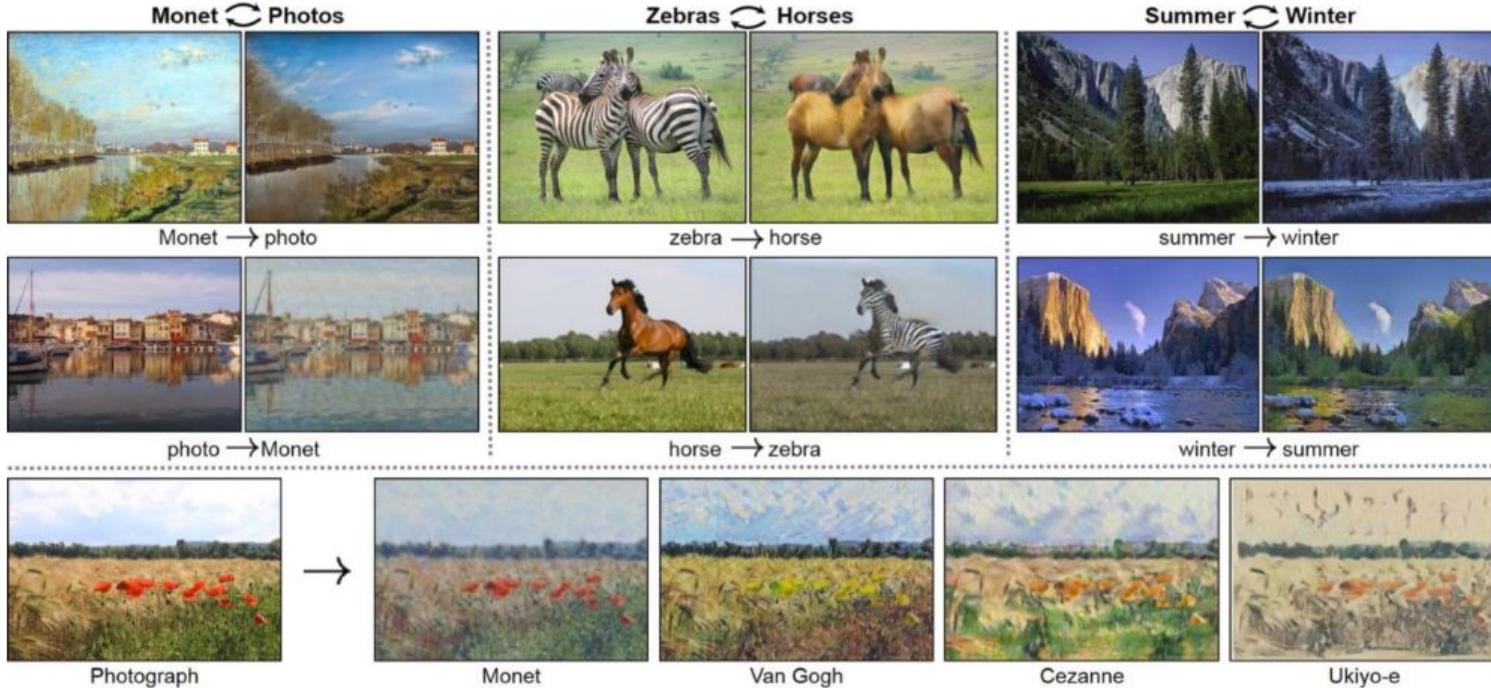


Image-to-Image Translation using GANs.

**Source:**

<https://medium.com/ai-society/gans-from-scratch-1-a-deep-introduction-with-code-in-pytorch-and-tensorflow-cb03cdcd8a0f>

# Art generation

*Christie's sold a portrait for \$432,000 that had been generated by a GAN, based on open-source code written by Robbie Barrat of Stanford. (He didn't see any of the money, it went to the French company)*

THE VERGE

REPORT

## HOW THREE FRENCH STUDENTS USED BORROWED CODE TO PUT THE FIRST AI PORTRAIT IN CHRISTIE'S

By James Vincent | Oct 23, 2018, 9:34am EDT



# (Recycle-)GAN: Fake video generation



# GAN Examples: Photo Generation

[thispersondoesnotexist.com](http://thispersondoesnotexist.com)

[thisxdoesnotexist.com](http://thisxdoesnotexist.com)



# GAN: Application Examples

---

- [AI Generates Fake Celebrity Faces \(Paper\)](#)
- [AI Learns Fashion Sense \(Paper\)](#)
- [Image to Image Translation using Cycle-Consistent Adversarial Neural Networks](#)
- [AI Creates Modern Art \(Paper\)](#)
- [This Deep Learning AI Generated Thousands of Creepy Cat Pictures](#)
- [MIT is using AI to create pure horror](#)
- [Amazon's new algorithm designs clothing by analyzing a bunch of pictures](#)
- [AI creates Photo-realistic Images \(Paper\)](#)

End

End

# GAN Examples: Music Generation

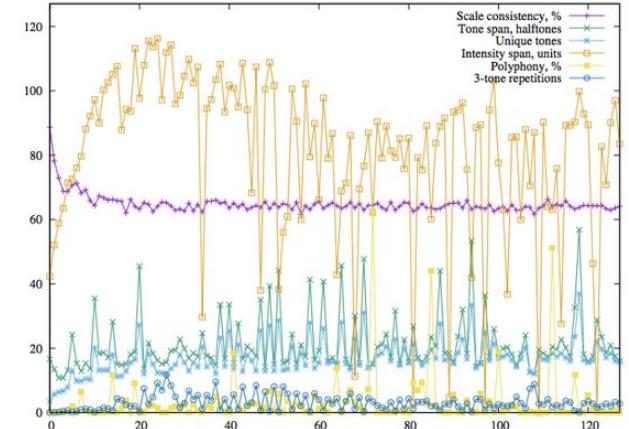
## C-RNN-GAN

*C-RNN\_GAN* is a recurrent neural network with adversarial training.

The adversaries are two different deep recurrent neural models, a generator ( $G$ ) and a discriminator ( $D$ ).

The generator is trained to generate data that is indistinguishable from real data, while the discriminator is trained to identify the generated data.

The training becomes a zero-sum game for which the Nash equilibrium is when the generator produces data that the discriminator cannot tell from real data.



(a) C-RNN-GAN using feature matching.

Source: <https://blog.sicara.com/keras-generative-adversarial-networks-image-deblurring-45e3ab6977b5>

# Neural Style Transfer

In order to obtain a description of an image's content, the image is fed forward through the CNN and network activations sampled at a late convolution layer of the VGG-19 architecture. Let us denote  $C(p)$  to indicate the content of the input image  $p$ .

In order to obtain a description of an image's style, the image is fed forward through the CNN and network activations sampled at early to middle layers of the VGG-19 architecture. These activations are encoded into a Gramian matrix representation. Let us denote  $S(a)$  to describe the style of the example style image  $a$ .

Image  $x$  is initially approximated by adding a small amount of white noise to input image  $p$ . An iterative optimization then proceeds, gradually updating  $x$  to minimize the error (loss function):

$$L(x) = |C(x) - C(p)| + k |S(x) - S(a)|$$

# GAN Basics: Generator Example Code

```
def generator(self):  
  
    model = Sequential()  
        model.add(Dense(256, input_shape=(100,)))  
        model.add(LeakyReLU(alpha=0.2))  
        model.add(BatchNormalization(momentum=0.8))  
        model.add(Dense(512))  
        model.add(LeakyReLU(alpha=0.2))  
        model.add(BatchNormalization(momentum=0.8))  
        model.add(Dense(1024))  
        model.add(LeakyReLU(alpha=0.2))  
        model.add(BatchNormalization(momentum=0.8))  
        model.add(Dense(self.WIDTH * self.HEIGHT *  
self.CHANNELS, activation='tanh'))  
  
    model.add(Reshape((self.WIDTH, self.HEIGHT, self.CHANNELS)))  
  
    return model
```

# GAN Basics: Discriminator Example Code

```
def discriminator(self):  
  
    model = Sequential()  
    model.add(Flatten(input_shape=self.SHAPE))  
    model.add(Dense((self.WIDTH * self.HEIGHT *  
self.CHANNELS), input_shape=self.SHAPE))  
  
    model.add(LeakyReLU(alpha=0.2))  
    model.add(Dense((self.WIDTH * self.HEIGHT *  
self.CHANNELS)/2))  
    model.add(LeakyReLU(alpha=0.2))  
  
    model.add(Dense(1, activation='sigmoid'))  
    model.summary()  
  
    return model
```

# Create the instances

```
def __init__(self, width = 28, height= 28, channels = 1):  
  
    self.WIDTH = width  
    self.HEIGHT = height  
    self.CHANNELS = channels  
  
    self.SHAPE = (self.WIDTH, self.HEIGHT, self.CHANNELS)  
  
    self.OPTIMIZER = Adam(lr=0.0002, decay=8e-9)  
  
    self.noise_gen = np.random.normal(0,1,(100,))  
  
    self.G = self.generator()  
    self.G.compile(loss='binary_crossentropy', \  
                   optimizer=self.OPTIMIZER)  
  
    self.D = self.discriminator()  
    self.D.compile(loss='binary_crossentropy', \  
                   optimizer=self.OPTIMIZER, metrics=['accuracy'])
```

# The Adversarial Model

*Just a Generator followed by a discriminator.*

*IMPORTANT: We froze the Discriminator weights.*

```
def stacked_G_D(self):
    self.D.trainable = False ## this freezes the weight, keep
                            ## reading to understand why this is
                            ## necessary

    model = Sequential()
    model.add(self.G)
    model.add(self.D)

    return model

### I need to compile the stacked model, so that a new instance of
### the discriminator is created with the frozen parameters.

self.stacked_G_D = self.stacked_G_D()
self.stacked_G_D.compile(loss='binary_crossentropy', \
optimizer=self.OPTIMIZER)
```

# Start Training (just batch training of discriminator)

```
def train(self, X_train, epochs=20000, batch = 32, save_interval = 200):

    for cnt in range(epochs):

        ## train discriminator

        random_index = np.random.randint(0, len(X_train) - batch/2)

        legit_images = X_train[random_index : random_index + batch/2].reshape(batch/2, self.WIDTH, self.HEIGHT, self.CHANNELS)

        gen_noise = np.random.normal(0, 1, (batch/2,100))
        syntetic_images = self.G.predict(gen_noise)

        x_combined_batch = np.concatenate((legit_images,
        syntetic_images))

        y_combined_batch = np.concatenate((np.ones((batch/2, 1)),
        np.zeros((batch/2, 1)))))

        d_loss = self.D.train_on_batch(x_combined_batch,
y_combined_batch)
```